Cortex RAG Test Document

Purpose

This sample PDF tests whether the ingestion pipeline can extract text from a PDF and make it searchable.

Sample facts

A RAG application retrieves relevant passages before asking a language model to generate an answer. The retrieved passages should provide evidence for the answer.

Snowflake Cortex Search can index text from a source query and return relevant records for a search query. The application can pass retrieved text to a language model and show the source filename to the user.

Questions to test

1. What does retrieval provide to a language model?
2. Why should a RAG app show source filenames?
3. What does the sample document say Cortex Search can do?